

# CSE 417: Artificial Intelligence

## Chapter 2: Intelligent Agents

Spring Semester 2015

Department of Computer Science and Engineering (CSE)


# Outline

- Agents and environments
- Rationality
- PEAS (Performance measure, Environment, Actuators, Sensors)
- Environment types
- Agent types

# Agents


- An **agent** is anything that can be viewed as **perceiving** its **environment** through **sensors** and **acting** upon that environment through **actuators**
- Human agent:
  - eyes, ears, and other organs for sensors;
  - hands, legs, mouth, and other body parts for actuators
- Robotic agent:
  - cameras and infrared range finders for sensors
  - various motors for actuators

# Agents and environments



- The **agent function** maps from percept histories to actions:  
 $[f: P^* \rightarrow A]$
- The **agent program** runs on the physical **architecture** to produce  $f$
- agent = architecture + program

# Vacuum-cleaner world



Demo:

<http://www.ai.sri.com/~oreilly/aima3java/aima3javademos.html>

- Percepts: location and contents, e.g., [A,Dirty]
- Actions: *Left*, *Right*, *Suck*, *NoOp*
- *Agent's function* → *look-up table*
  - *For many agents this is a very large table*

Percept sequence	Action
[A, Clean]	<i>Right</i>
[A, Dirty]	<i>Suck</i>
[B, Clean]	<i>Left</i>
[B, Dirty]	<i>Suck</i>
[A, Clean], [A, Clean]	<i>Right</i>
[A, Clean], [A, Dirty]	<i>Suck</i>
⋮	⋮

# Rational agents

- **Rationality**
  - Performance measuring success
  - Agents prior knowledge of environment
  - Actions that agent can perform
  - Agent's percept sequence to date
- **Rational Agent:** For each possible percept sequence, a rational agent should select an action that is expected to maximize its performance measure, given the evidence provided by the percept sequence and whatever built-in knowledge the agent has.

•

# Rationality

- Rational is different from omniscience
  - Percepts may not supply all relevant information
  - E.g., in card game, don't know cards of others.
- Rational is different from being perfect
  - Rationality maximizes expected outcome while perfection maximizes actual outcome.

# Autonomy in Agents

**The autonomy of an agent is the extent to which its behaviour is determined by its own experience, rather than knowledge of designer.**

- Extremes
  - No autonomy – ignores environment/data
  - Complete autonomy – must act randomly/no program
- Example: baby learning to crawl
- Ideal: design agents to have some autonomy
  - Possibly become more autonomous with experience



# PEAS

- PEAS: Performance measure, Environment, Actuators, Sensors
- Must first specify the setting for intelligent agent design
- Consider, e.g., the task of designing an automated taxi driver:
  - Performance measure: Safe, fast, legal, comfortable trip, maximize profits
  - Environment: Roads, other traffic, pedestrians, customers
  - Actuators: Steering wheel, accelerator, brake, signal, horn
  - Sensors: Cameras, sonar, speedometer, GPS, odometer, engine sensors, keyboard

# PEAS

- Agent: Part-picking robot
- Performance measure: Percentage of parts in correct bins
- Environment: Conveyor belt with parts, bins
- Actuators: Jointed arm and hand
- Sensors: Camera, joint angle sensors

# PEAS

- Agent: Interactive English tutor
- Performance measure: Maximize student's score on test
- Environment: Set of students
- Actuators: Screen display (exercises, suggestions, corrections)
- Sensors: Keyboard

# Environment types

- Fully observable (vs. partially observable)
- Deterministic (vs. stochastic)
- Episodic (vs. sequential)
- Static (vs. dynamic)
- Discrete (vs. continuous)
- Single agent (vs. multiagent):

## Fully observable (vs. partially observable)

- Is everything an agent requires to choose its actions available to it via its sensors? Perfect or Full information.
  - If so, the environment is fully accessible
- If not, parts of the environment are inaccessible
  - Agent must make informed guesses about world.
- In decision theory: perfect information vs. imperfect information.

**Cross Word**  
Fully

**Poker**  
Partially

**Backgammon**  
Partially

**Taxi driver**  
Partially

**Part picking robot**  
Fully

**Image analysis**  
Fully

## Deterministic (vs. stochastic)

- Does the change in world state
  - Depend only on current state and agent's action?
- Non-deterministic environments
  - Have aspects beyond the control of the agent
  - Utility functions have to guess at changes in world

<b>Cross Word</b>	<b>Poker</b>	<b>Backgammon</b>	<b>Taxi driver</b>	<b>Part picking robot</b>	<b>Image analysis</b>
Deterministic	Stochastic	Stochastic	Stochastic	Stochastic	Deterministic

## Episodic (vs. sequential):

- Is the choice of current action
  - Dependent on previous actions?
  - If not, then the environment is episodic
- In non-episodic environments:
  - Agent has to plan ahead:
    - Current choice will affect future actions

<b>Cross Word</b>	<b>Poker</b>	<b>Backgammon</b>	<b>Taxi driver</b>	<b>Part picking robot</b>	<b>Image analysis</b>
Sequential	Sequential	Sequential	Sequential	Episodic	Episodic

## Static (vs. dynamic):

- Static environments don't change
  - While the agent is deliberating over what to do
- Dynamic environments do change
  - So agent should/could consult the world when choosing actions
  - Alternatively: anticipate the change during deliberation OR make decision very fast
- Semidynamic: If the environment itself does not change with the passage of time but the agent's performance score does.

<b>Cross Word</b>	<b>Poker</b>	<b>Backgammon</b>	<b>Taxi driver</b>	<b>Part picking robot</b>	<b>Image analysis</b>
Static	Static	Static	Dynamic	Dynamic	Semi

Another example: off-line route planning vs. on-board navigation system



## Discrete (vs. continuous)

- A limited number of distinct, clearly defined percepts and actions vs. a range of values (continuous)

<b>Cross Word</b>	<b>Poker</b>	<b>Backgammon</b>	<b>Taxi driver</b>	<b>Part picking robot</b>	<b>Image analysis</b>
Discrete	Discrete	Discrete	Conti	Conti	Conti

## Single agent (vs. multiagent):

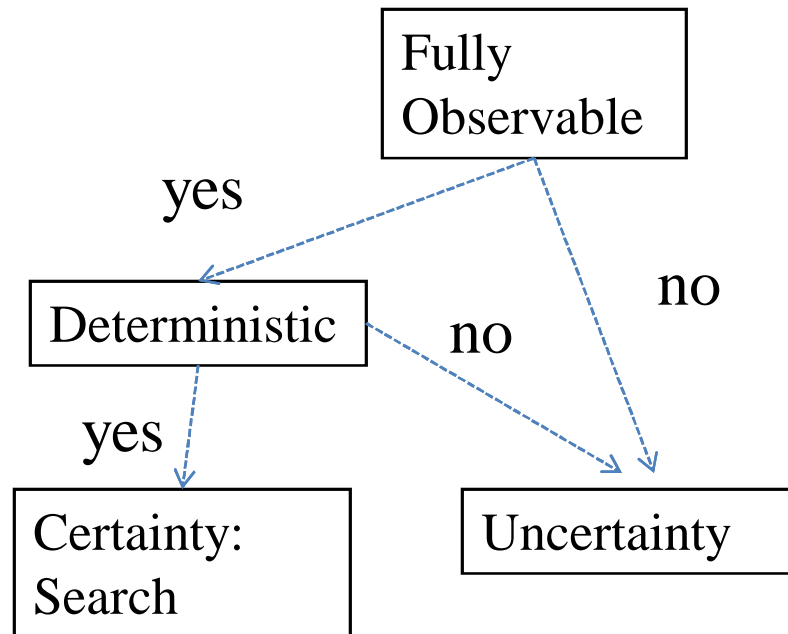
- An agent operating by itself in an environment or there are many agents working together

<b>Cross Word</b>	<b>Poker</b>	<b>Backgammon</b>	<b>Taxi driver</b>	<b>Part picking robot</b>	<b>Image analysis</b>
Single	Multi	Multi	Multi	Single	Single

# Summary.

	Observable	Deterministic	Episodic	Static	Discrete	Agents
<b>Cross Word</b>	Fully	Deterministic	Sequential	Static	Discrete	Single
<b>Poker</b>	Fully	Stochastic	Sequential	Static	Discrete	Multi
<b>Backgammon</b>	Partially	Stochastic	Sequential	Static	Discrete	Multi
<b>Taxi driver</b>	Partially	Stochastic	Sequential	Dynamic	Conti	Multi
<b>Part picking robot</b>	Partially	Stochastic	Episodic	Dynamic	Conti	Single
<b>Image analysis</b>	Fully	Deterministic	Episodic	Semi	Conti	Single

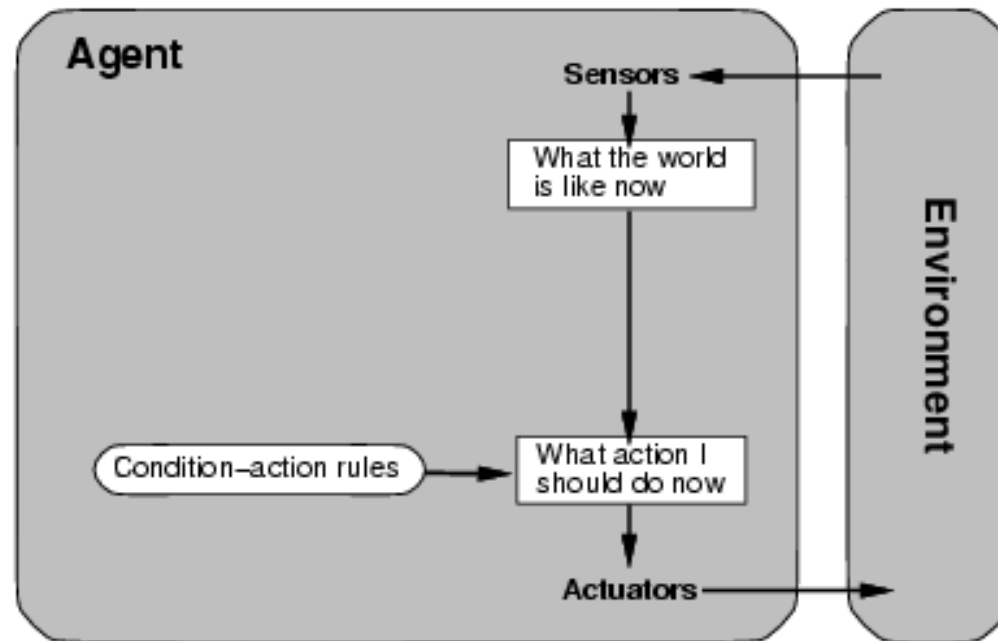
# Choice under (Un)certainty



# Agent types

- Four basic types in order of increasing generality:
  - Simple reflex agents
  - Reflex agents with state/model
  - Goal-based agents
  - Utility-based agents
  - All these can be turned into learning agents
  - <http://www.ai.sri.com/~oreilly/aima3ejava/aima3ejavademos.html>

# Simple reflex agents



```
function REFLEX-VACUUM-AGENT( [location,status] ) returns an action
  if status = Dirty then return Suck
  else if location = A then return Right
  else if location = B then return Left
```

# Simple reflex agents

- Simple but very limited intelligence.
- **Action does not depend on percept history, only on current percept.**
- Therefore no memory requirements.
- Infinite loops
  - Suppose vacuum cleaner does not observe location. What do you do given location = clean? Left of A or right on B -> infinite loop.
  - [Fly buzzing](#) around window or light.
  - Possible Solution: Randomize action.
  - Thermostat.
- Chess – openings, endings
  - Lookup table (not a good idea in general)
    - $35^{100}$  entries required for the entire game

# States: Beyond Reflexes

- Recall the **agent function** that maps from percept histories to actions:

$$[f: P^* \rightarrow A]$$

- An agent program can implement an agent function by maintaining an **internal state**.
- The internal state can contain information about the state of the external environment.
- The state depends on the history of percepts and on the history of actions taken:

$$[f: P^*, A^* \rightarrow S \rightarrow A] \text{ where } S \text{ is the set of states.}$$

- If each internal state includes all information relevant to information making, the state space is **Markovian**.



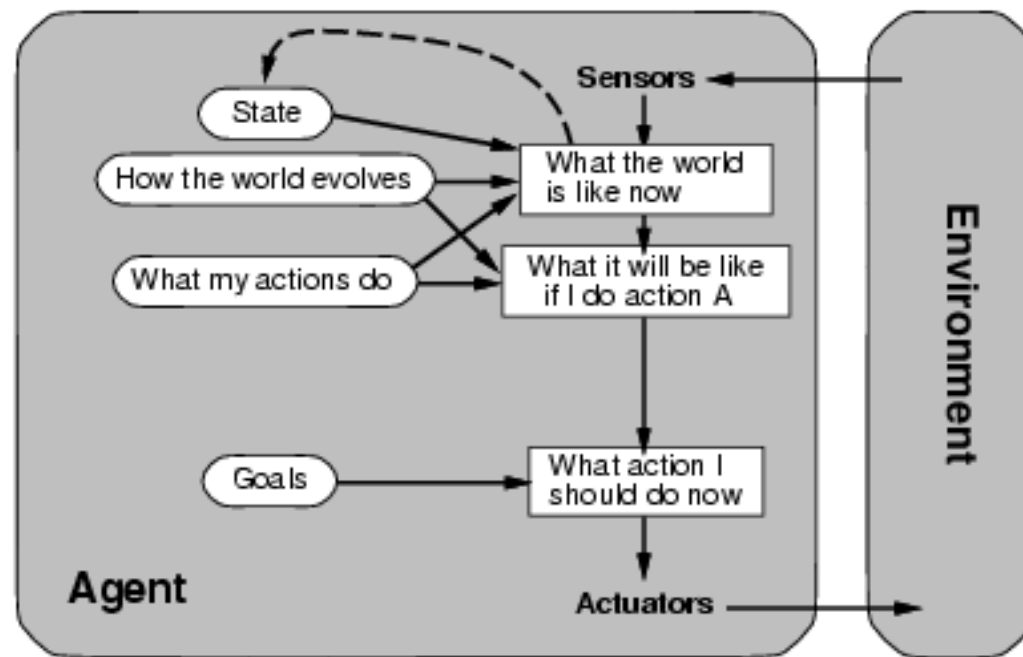
# States and Memory: Game Theory

- If each state includes the information about the percepts and actions that led to it, the state space has **perfect recall**.
- **Perfect Information** = Perfect Recall + Full Observability.

# Goal-based agents

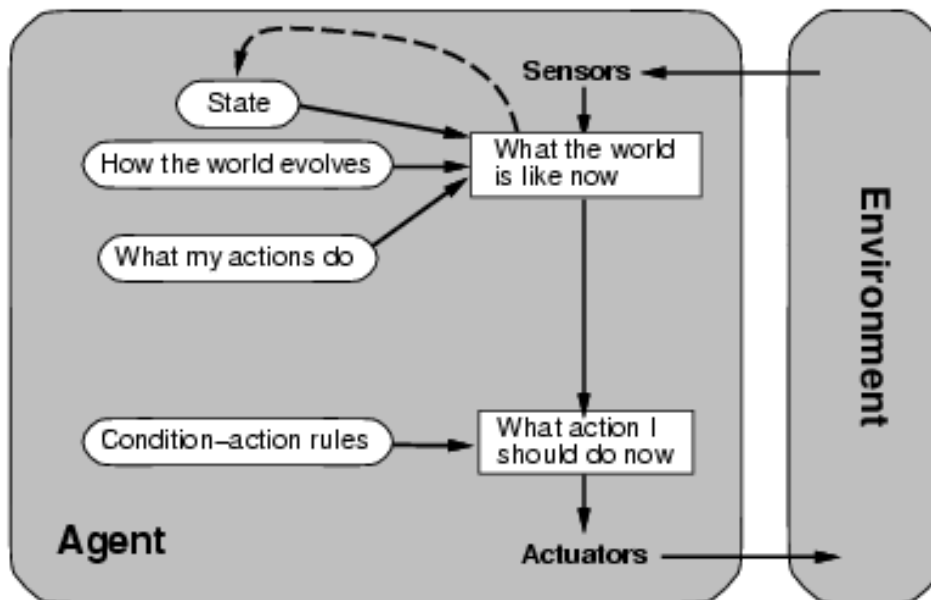
- knowing state and environment? Enough?
  - Taxi can go left, right, straight
- Have a goal
  - A destination to get to
- Uses knowledge about a goal to guide its actions
  - E.g., Search, planning

# Goal-based agents



- Reflex agent breaks when it sees brake lights. Goal based agent reasons
  - Brake light -> car in front is stopping -> I should stop -> I should use brake

# Model-based reflex agents



- Know how world evolves
  - Overtaking car gets closer from behind
- How agents actions affect the world
  - Wheel turned clockwise takes you right

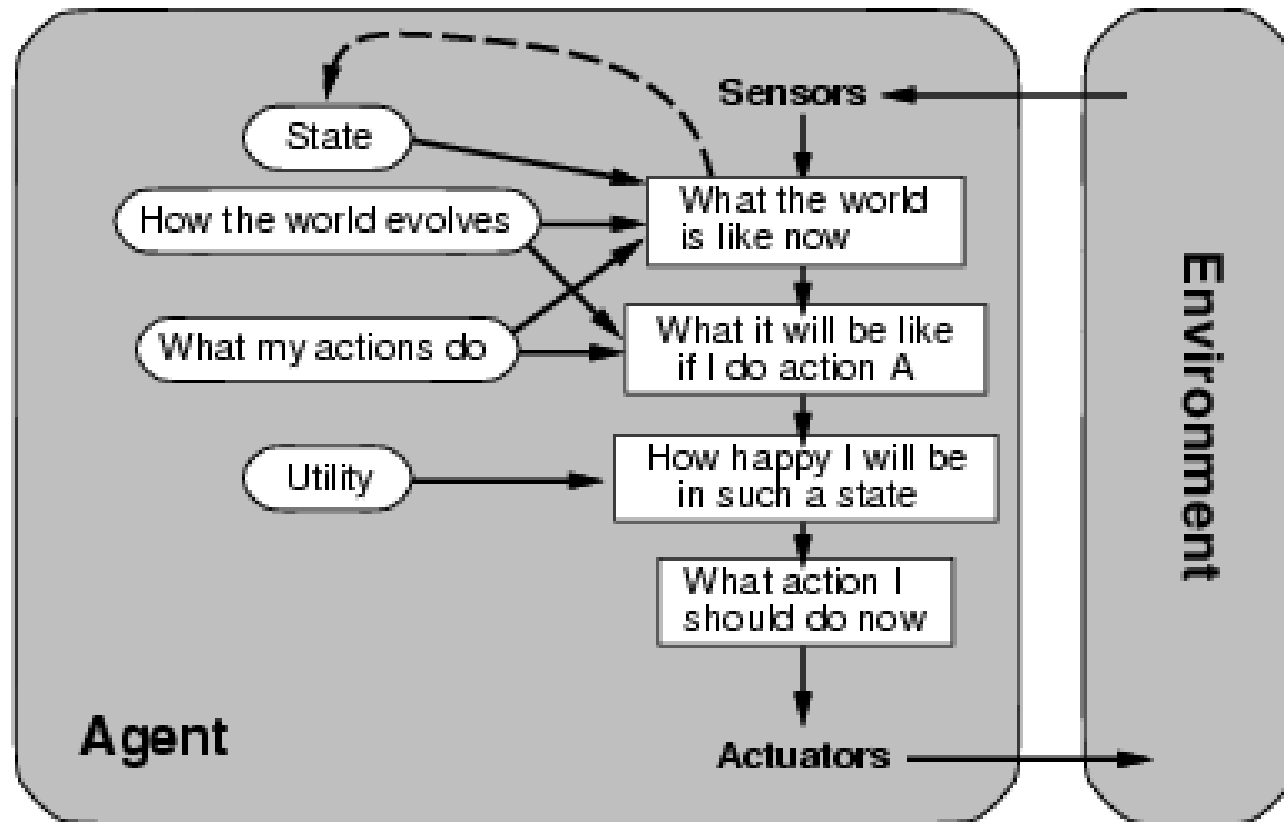
```
function REFLEX-AGENT-WITH-STATE(percept) returns action
  static: state, a description of the current world state
         rules, a set of condition-action rules

  state ← UPDATE-STATE(state, percept)
  rule ← RULE-MATCH(state, rules)
  action ← RULE-ACTION[rule]
  state ← UPDATE-STATE(state, action)
  return action
```

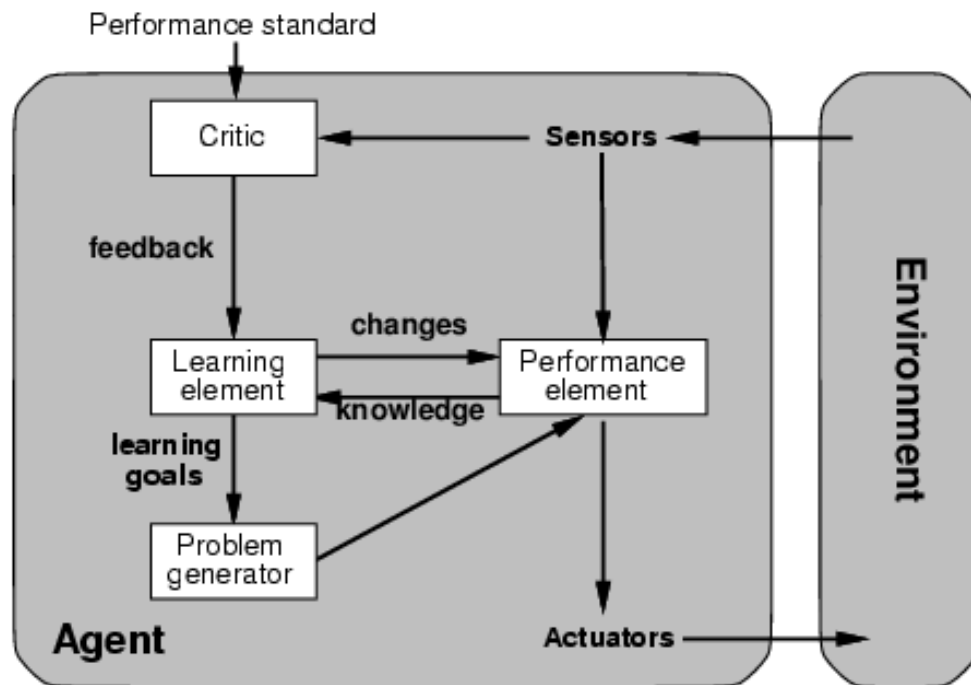
# Utility-based agents

- Goals are not always enough
  - Many action sequences get taxi to destination
  - Consider other things. How fast, how safe.....
- A utility function maps a state onto a real number which describes the associated degree of “happiness”, “goodness”, “success”.
- Where does the utility measure come from?
  - Economics: money.
  - Biology: number of offspring.
  - Your life?

# Utility-based agents

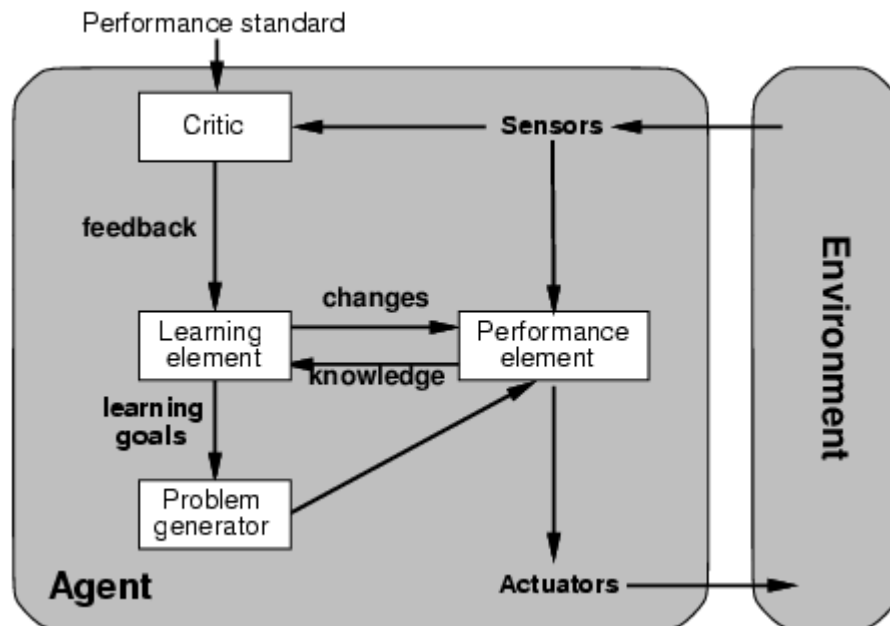


# Learning agents



- Performance element is what was previously the whole agent
  - Input sensor
  - Output action
- Learning element
  - Modifies performance element.

# Learning agents



- Critic: how the agent is doing
  - Input: checkmate?
  - Fixed
- Problem generator
  - Tries to solve the problem differently instead of optimizing.
  - Suggests **exploring** new actions -> new problems.



# Learning agents(Taxi driver)

- Performance element
  - How it currently drives
- Taxi driver Makes quick left turn across 3 lanes
  - Critics observe shocking language by passenger and other drivers and informs bad action
  - Learning element tries to modify performance elements for future
  - Problem generator suggests experiment out something called Brakes on different Road conditions
- Exploration vs. Exploitation
  - Learning experience can be costly in the short run
  - shocking language from other drivers
  - Less tip
  - Fewer passengers