

# CSE 517: Artificial Intelligence

## Chapter 3: Search and Sequential Action

Spring 2015

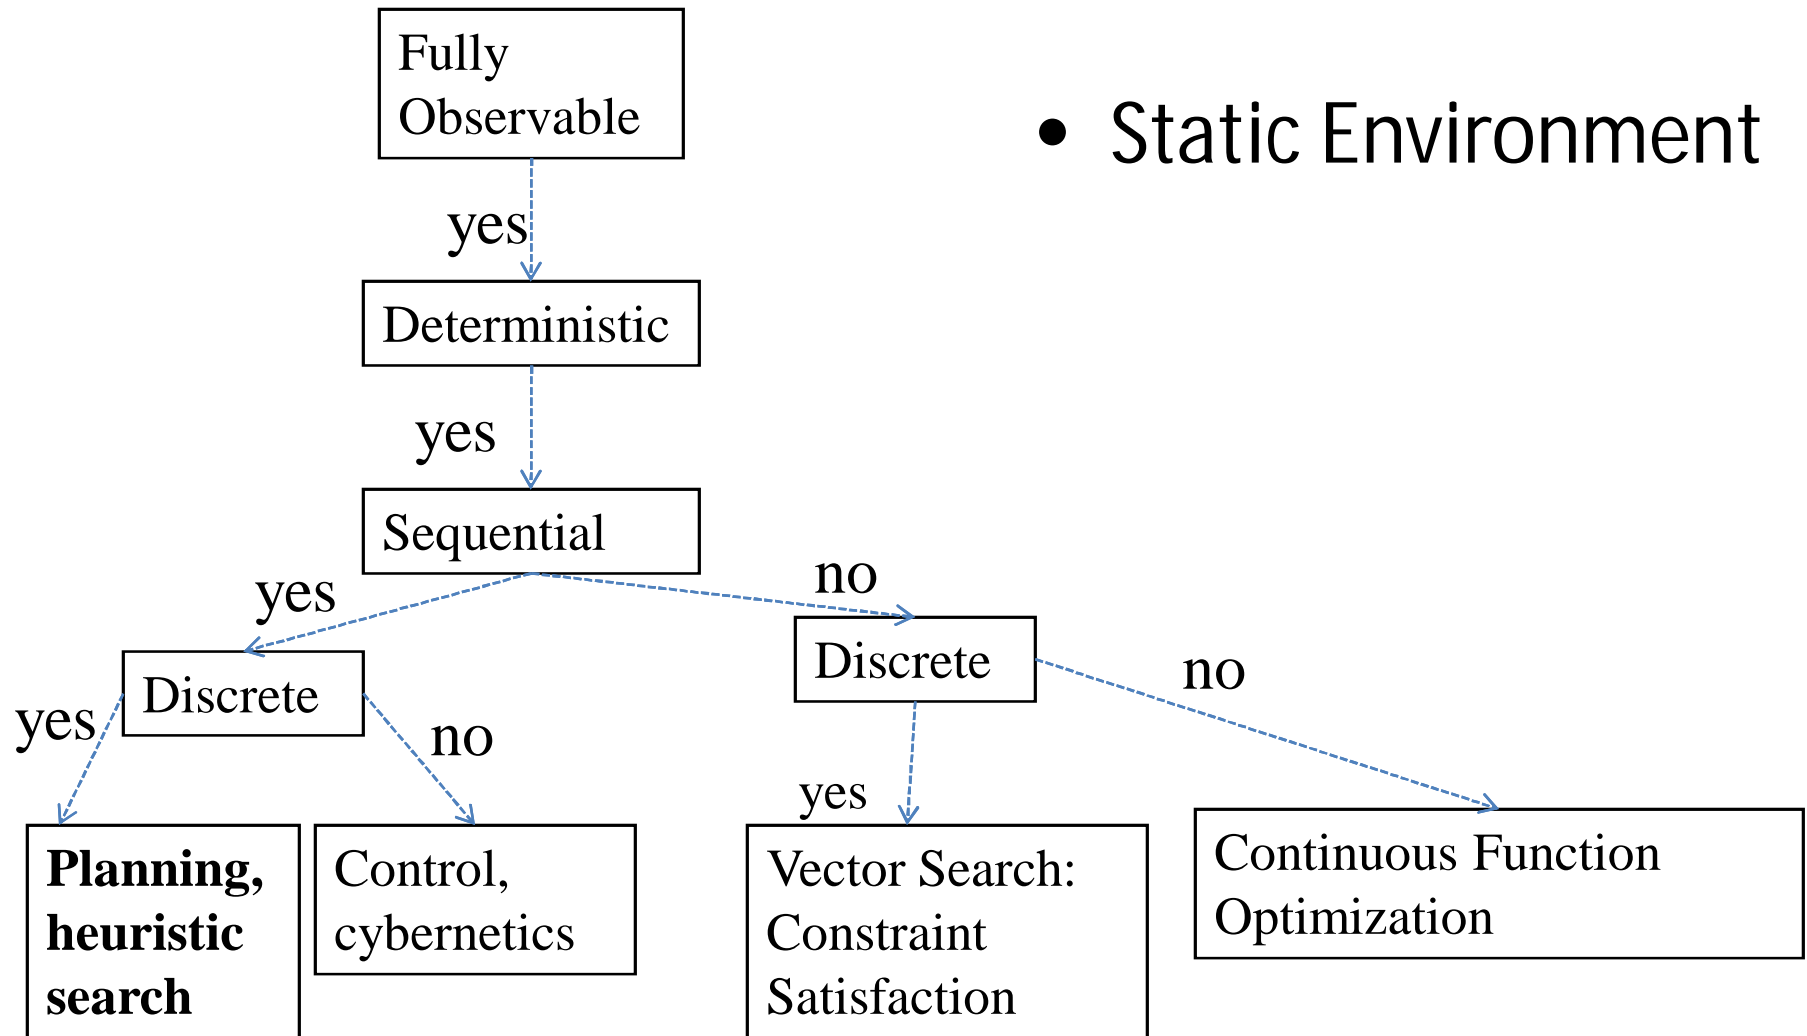
Department of Computer Science and Engineering (CSE)

# Outline

- Problem formulation: representing sequential problems.
- Example problems.
- Planning for solving sequential problems without uncertainty.
- Basic search algorithms

# Environment Type Discussed In this Lecture

- Static Environment



## Choice in a Deterministic Known Environment

- Without uncertainty, choice is trivial in principle: choose what you know to be the best option.
- Trivial if the problem is represented in a look-up table.

Option	Value
Chocolate	10
Wine	20
Book	15

This is the standard problem representation in decision theory (economics).

# Computational Choice Under Certainty


- But choice can be *computationally* hard if the problem information is represented differently.
- Options may be **structured** and the best option needs to be constructed.
  - E.g., an option may consist of a path, sequence of actions, plan, or strategy.
- The value of options may be given **implicitly** rather than explicitly.
  - E.g., cost of paths need to be computed from map.

# Sequential Action Example

- Deterministic, fully observable → single-state problem
  - Agent knows exactly which state it will be in; solution is a sequence
  - Vacuum world → everything observed
  - Romania → The full map is observed


- Single-state: Start in #5. Solution??
  - [Right, Suck]

•



# Problem types

- Non-observable → sensorless problem (conformant problem)
  - Agent may have no idea where it is; solution is a sequence
  - Vacuum world → No sensors
  - Romania → No map just know operators(cities you can move to)
- Conformant: Start in {1, 2, 3, 4, 5, 6, 7, 8}
  - e.g., Right goes to {2, 4, 6, 8}. Solution??
  - [Right, Suck, Left, Suck]




# Problem types

- Nondeterministic and/or partially observable → contingency problem

- percepts provide new information about current state
- Unknown state space → exploration problem
- Vacuum world → know state of current location
- Romania → know current location and neighbor cities

- Contingency: [L,clean]

- Start in #5 or #7
- Murphy's Law: Suck can dirty a clean carpet
- Local sensing: dirt, location only.
- Solution??
- [Right, if dirt then Suck]

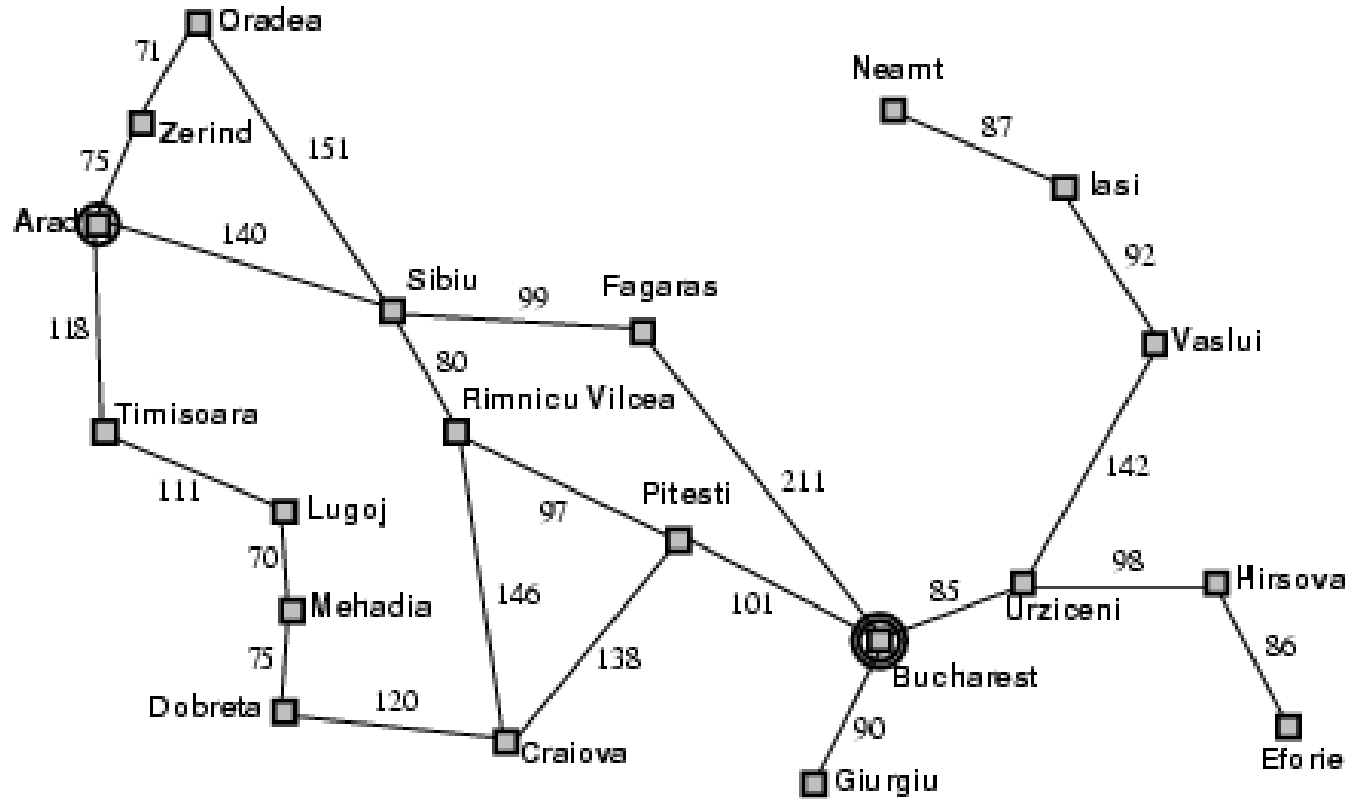




# Example: Romania

- On holiday in Romania; currently in Arad.
- Flight leaves tomorrow from Bucharest
- 
- Formulate goal:
  - be in Bucharest
  -
- Formulate problem:
  - **states**: various cities
  - **actions**: drive between cities
- Find solution:
  - sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest
  -

# Example: Romania



Abstraction: The process of removing details from a representation  
Is the map a good representation of the problem? What is a good replacement?

# Single-state problem formulation


A **problem** is defined by four items:

1. **initial state** e.g., "at Arad"
  2. **actions** or **successor function**  $S(x)$  = set of action–state pairs
    - e.g.,  $S(Arad) = \{ \langle Arad \rightarrow Zerind, Zerind \rangle, \dots \}$
    -
  3. **goal test**, can be
    - **explicit**, e.g.,  $x = \text{"at Bucharest"}$
    - **implicit**, e.g.,  $Checkmate(x)$
    -
  4. **path cost** (additive)
    - e.g., sum of distances, number of actions executed, etc.
    - $c(x,a,y)$  is the **step cost**, assumed to be  $\geq 0$
    -
- A **solution** is
    - a sequence of actions leading from the initial state to a goal state
    -

# Selecting a state space


- Real world is complex
  - state space must be **abstracted** for problem solving
- (Abstract) state = set of real states
- 
- (Abstract) action = complex combination of real actions
  - e.g., "Arad → Zerind" represents a complex set of possible routes, detours, rest stops, etc.
- (Abstract) solution =
  - set of real paths that are solutions in the real world
  -
- Each abstract action should be "easier" than the original problem
-

# Vacuum world state space graph



- states?
- actions?
- goal test?
- path cost?
-

# Vacuum world state space graph



- states? integer dirt and robot location
- actions? *Left, Right, Suck*
- goal test? no dirt at all locations
- path cost? 1 per action

# Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- states?
- actions?
- goal test?
- path cost?

# Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8


Goal State

- states? locations of tiles
- actions? move blank left, right, up, down
- goal test? = goal state (given)
- path cost? 1 per move

[Note: optimal solution of  $n$ -Puzzle family is NP-hard]



# Example: robotic assembly



- states?:
  - real-valued coordinates of robot joint angles
  - parts of the object to be assembled
  -
- actions?:
  - continuous motions of robot joints
- goal test?:
  - complete assembly
  -
- path cost?:
  - time to execute
  -

# Problem-solving agents

```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  static: seq, an action sequence, initially empty
          state, some description of the current world state
          goal, a goal, initially null
          problem, a problem formulation

  state ← UPDATE-STATE(state, percept)
  if seq is empty then
    goal ← FORMULATE-GOAL(state)
    problem ← FORMULATE-PROBLEM(state, goal)
    seq ← SEARCH(problem)
  action ← FIRST(seq)
  seq ← REST(seq)
  return action
```

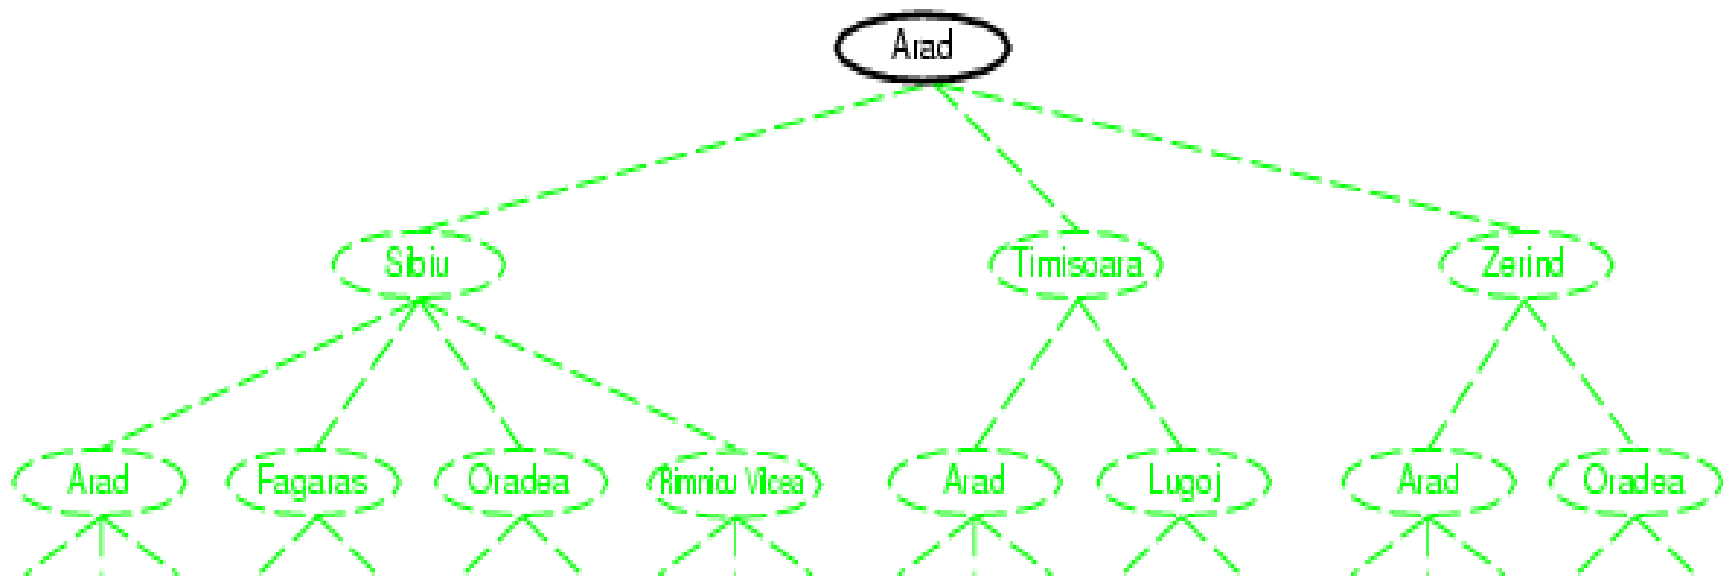
Note: this is offline problem solving; solution executed “eyes closed.”

# Tree search algorithms

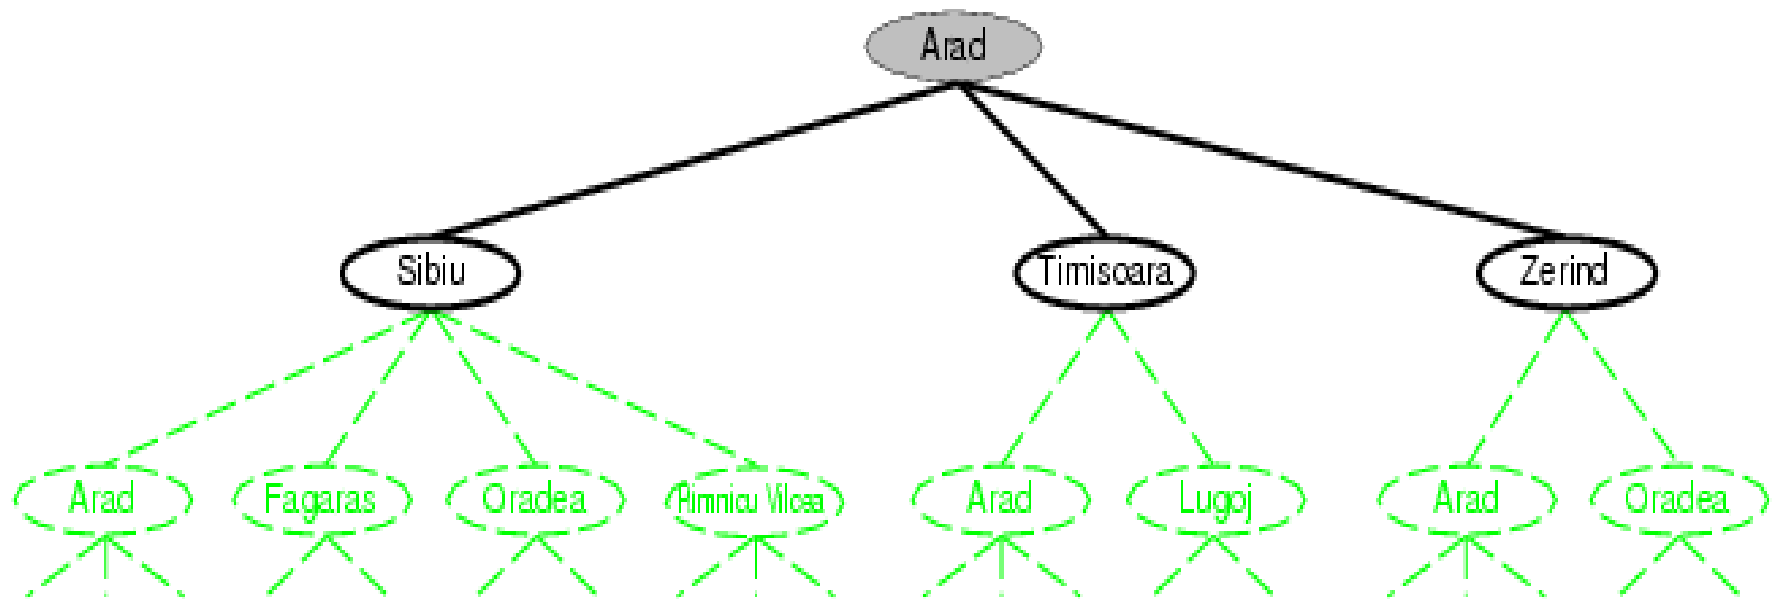
- Basic idea:
  - offline, simulated exploration of state space by generating successors of already-explored states (a.k.a. ~expanding states)

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
```

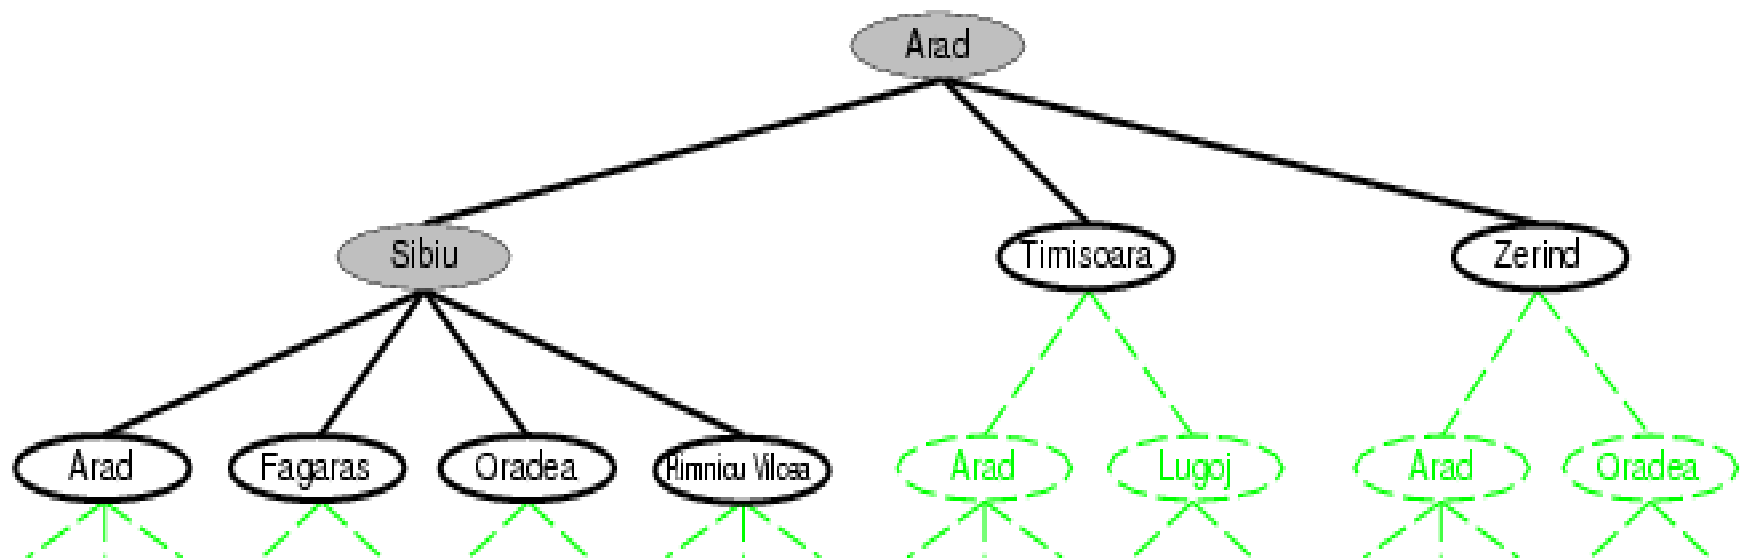
# Tree search example



# Tree search example



# Tree search example



# Search Graph vs. State Graph

- Be careful to distinguish
  - Search tree: nodes are **sequences of actions**.
  - State Graph: Nodes are states of the environment.
  - We will also consider soon **search graphs**.
- Demo: <http://aispace.org/search/>

# Search strategies

- A search strategy is defined by picking the **order of node expansion**
- Strategies are evaluated along the following dimensions:
  - **completeness**: does it always find a solution if one exists?
  - **time complexity**: number of nodes generated
  - **space complexity**: maximum number of nodes in memory
  - **optimality**: does it always find a least-cost solution?
  -
- Time and space complexity are measured in terms of
  - $b$ : maximum branching factor of the search tree
  - $d$ : depth of the least-cost solution
  - $m$ : maximum depth of the state space (may be  $\infty$ )
  -




# Uninformed search strategies

- **Uninformed** search strategies use only the information available in the problem definition
- Breadth-first search
- Depth-first search
- Depth-limited search
- Iterative deepening search

# Breadth-first search


- Expand shallowest unexpanded node
- 
- **Implementation:**
  - *Frontier* or *fringe* is a FIFO queue, i.e., new successors go at the end



# Breadth-first search

- Expand shallowest unexpanded node
- 
- **Implementation:**
  - *frontier* is a FIFO queue, i.e., new successors go at end


–



# Breadth-first search

- Expand shallowest unexpanded node
- 
- **Implementation:**
  - *frontier* is a FIFO queue, i.e., new successors go at end


–



# Breadth-first search

- Expand shallowest unexpanded node
- <http://aispace.org/search/>
- **Implementation:**
  - *frontier* is a FIFO queue, i.e., new successors go at end

–




# Properties of breadth-first search

- Complete? Time? Space? Optimal?
- Complete? Yes (if  $b$  is finite)
- 
- Time?  $1+b+b^2+b^3+\dots +b^d + b(b^d-1) = O(b^{d+1})$
- 
- Space?  $O(b^{d+1})$  (keeps every node in memory)
- 
- Optimal? Yes (if cost = 1 per step)
- 
- **Space** is the bigger problem (more than time)
-


# Depth-first search

- Expand deepest unexpanded node
- 
- **Implementation:**
  - *frontier* = LIFO queue, i.e., put successors at front
  -



# Depth-first search


- Expand deepest unexpanded node
- 
- **Implementation:**
  - *frontier* = LIFO queue, i.e., put successors at front
  -






# Depth-first search

- Expand deepest unexpanded node
- 
- **Implementation:**
  - *frontier* = LIFO queue, i.e., put successors at front
  -




# Depth-first search

- Expand deepest unexpanded node
- 
- **Implementation:**
  - *frontier* = LIFO queue, i.e., put successors at front
  -




# Depth-first search

- Expand deepest unexpanded node
- 
- **Implementation:**
  - *frontier* = LIFO queue, i.e., put successors at front
  -




# Depth-first search

- Expand deepest unexpanded node
- 
- **Implementation:**
  - *frontier* = LIFO queue, i.e., put successors at front
  -




# Depth-first search

- Expand deepest unexpanded node
- 
- **Implementation:**
  - *frontier* = LIFO queue, i.e., put successors at front
  -




# Depth-first search

- Expand deepest unexpanded node
- 
- **Implementation:**
  - *frontier* = LIFO queue, i.e., put successors at front
  -




# Depth-first search

- Expand deepest unexpanded node
- 
- **Implementation:**
  - *frontier* = LIFO queue, i.e., put successors at front
  -



# Depth-first search


- Expand deepest unexpanded node
- 
- **Implementation:**
  - *frontier* = LIFO queue, i.e., put successors at front
  -






# Depth-first search

- Expand deepest unexpanded node
- 
- **Implementation:**
  - *frontier* = LIFO queue, i.e., put successors at front
  -



# Depth-first search

- Expand deepest unexpanded node
- <http://aispace.org/search/>
- **Implementation:**
  - *frontier* = LIFO queue, i.e., put successors at front
  -



# Properties of depth-first search

- Complete? Time? Space? Optimal?
- Complete? No: fails in infinite-depth spaces, spaces with loops
  - Modify to avoid repeated states along path (graph search)
  - - complete in finite spaces
- Time?  $O(b^m)$ : terrible if maximum depth  $m$  is much larger than solution depth  $d$ 
  - but if solutions are dense, may be much faster than breadth-first
  -
- Space?  $O(bm)$ , i.e., linear space! Store single path with unexpanded siblings.
  - Seems to be common in animals and humans.
  -
- Optimal? No.
- Important for exploration (on-line search).
-

# Depth-limited search

- depth-first search with depth limit  $l$ ,
  - i.e., nodes at depth  $l$  have no successors
  - Solves infinite loop problem
- Common AI strategy: let user choose search/resource bound.
- Complete? No if  $l < d$ :
- Time?  $O(b^l)$ :
- Space?  $O(bl)$ , i.e., linear space!
- Optimal? No if  $l > b$

# Iterative deepening search

**function** ITERATIVE-DEEPENING-SEARCH(*problem*) **returns** a solution, or failure

**inputs:** *problem*, a problem

**for** *depth*  $\leftarrow 0$  **to**  $\infty$  **do**

*result*  $\leftarrow$  DEPTH-LIMITED-SEARCH(*problem*, *depth*)

**if** *result*  $\neq$  cutoff **then return** *result*


# Iterative deepening search $\neq 0$

Limit = 0




# Iterative deepening search $l=1$

Limit = 1



# Iterative deepening search $l=2$


Limit = 2





# Iterative deepening search $l=3$

Limit = 3



# Iterative deepening search

- Number of nodes generated in a depth-limited search to depth  $d$  with branching factor  $b$ :

$$N_{DLS} = b^0 + b^1 + b^2 + \dots + b^{d-2} + b^{d-1} + b^d$$

- Number of nodes generated in an iterative deepening search to depth  $d$  with branching factor  $b$ :

$$N_{IDS} = (d+1)b^0 + d b^1 + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d$$

- For  $b = 10$ ,  $d = 5$ ,

- 

- $N_{DLS} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$

- 

- $N_{IDS} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$

- 

- Overhead =  $(123,456 - 111,111)/111,111 = 11\%$

# Properties of iterative deepening search

- Complete? Yes
- 
- Time?  $(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- 
- Space?  $O(bd)$
- 
- Optimal? Yes, if step cost = 1



# Summary of algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Time	$b^d$	$b^d$	$b^m$	$b^l$	$b^d$	$b^{d/2}$
Space	$b^d$	$b^d$	$bm$	$bl$	$bd$	$b^{d/2}$
Optimal?	Yes	Yes	No	No	Yes	Yes
Complete?	Yes	Yes	No	Yes, if $l \geq d$	Yes	Yes

# Repeated states

- Failure to detect repeated states can turn a linear problem into an exponential one!




# Graph search


```
function GRAPH-SEARCH(problem, fringe) returns a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    if STATE[node] is not in closed then
      add STATE[node] to closed
      fringe ← INSERTALL(EXPAND(node, problem), fringe)
```

- Simple solution: just keep track of which states you have visited.
- Usually easy to implement in modern computers.


# The Separation Property of Graph Search



(a)



(b)



(c)

- Black: expanded nodes.
- White: frontier nodes.
- Grey: unexplored nodes.

# Summary

- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored
- 
- Variety of uninformed search strategies
- 
- Iterative deepening search uses only linear space and not much more time than other uninformed algorithms
-